

# CPU SCHEDULING

Shivaram Venkataraman  
CS 537, Spring 2019



OK?

# ADMINISTRIVIA

- Project Ia is due ~~today!~~ **Thursday at 11.59pm**
- No office hours from 5pm Tue to noon Thu
- Fill out office hours form? <https://goo.gl/forms/5VxrwRawtEFkrjO23>
- No more waitlist!
- Project Ib out tomorrow. Schedule updates

# AGENDA / LEARNING OUTCOMES

## Scheduling

How does the OS decide what process to run?

What are some of the metrics to optimize for?

## Policies

How to handle interactive and batch processes?

What to do when OS doesn't have complete information?

**RECAP**

# RECAP: SCHEDULING MECHANISM

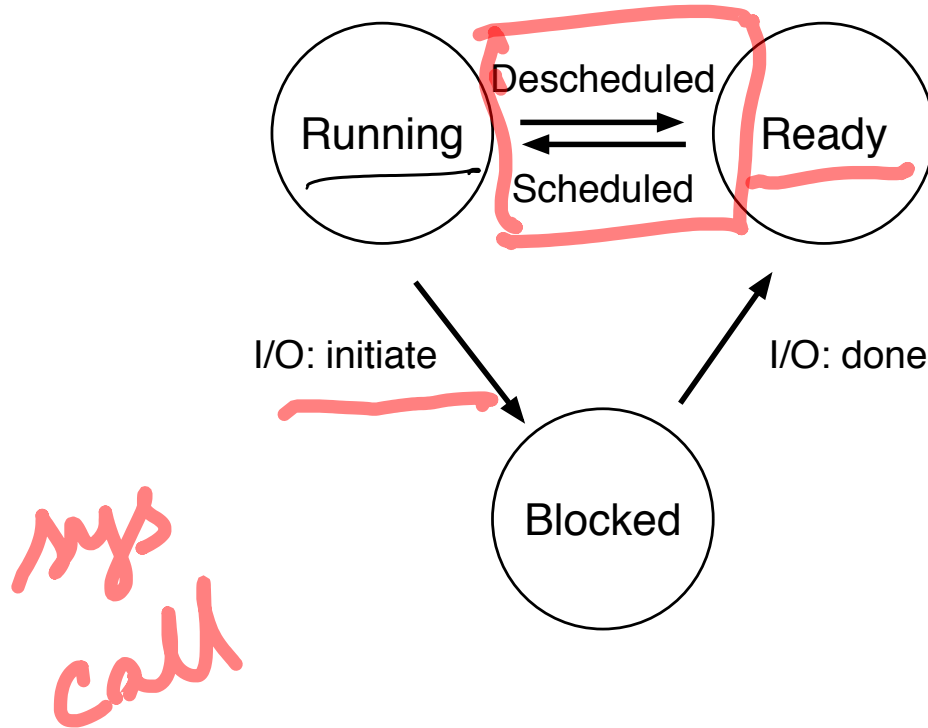
Process: Abstraction to virtualize CPU

Address Space  
File Desc.

Use time-sharing in OS to switch between processes

P1 → pause  
[CPU] ← P2

# PROCESS STATE TRANSITIONS



# RECAP: SCHEDULING MECHANISM

Limited Direct Execution



Use system calls to run access devices etc. from user mode

Context-switch using interrupts for multi-tasking

# Operating System

# Hardware

# Program Process A

Cat  
my file

timer interrupt  
save regs(A) to k-stack(A)  
move to kernel mode  
jump to trap handler

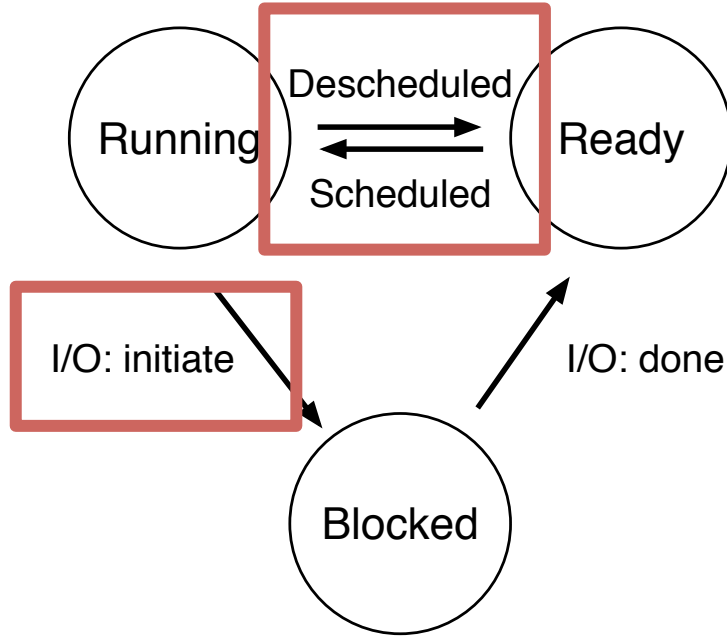
Handle the trap  
Call switch() routine  
save kernel regs(A) to proc-struct(A)  
restore kernel regs(B) from proc-struct(B)  
switch to k-stack(B)  
return-from-trap (into B)

user stack pointer

restore regs(B) from k-stack(B)  
move to user mode  
jump to B's IP

Process B





POLICY ?

# VOCABULARY

Workload: set of **jobs** (arrival time, run\_time)

**Job** ~ Current execution of a process

Alternates between CPU and I/O

Moves between ready and blocked queues

**Scheduler**: Decides which ready job to run

Metric: measurement of scheduling quality

$J_1$

$t = 0$

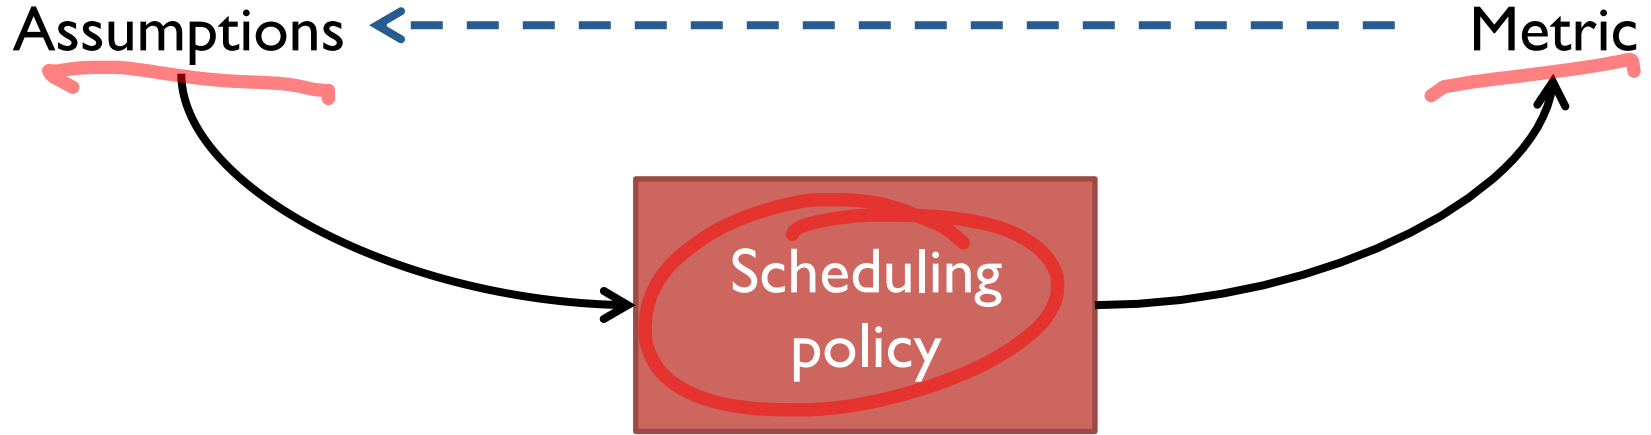
$r = 10s$

$J_2$

$t = 10$

$r = 20s$

# APPROACH



# ASSUMPTIONS

1. Each job runs for the same amount of time
2. All jobs arrive at the same time
3. All jobs only use the CPU (no I/O)
4. Run-time of each job is known

# METRIC 1: TURNAROUND TIME

Turnaround time = *completion\_time* - *arrival\_time*

Example:

Process A arrives at time  $t = 10$ , finishes  $t = 30$

Process B arrives at time  $t = 10$ , finishes  $t = 50$

Turnaround time

A = 20, B = 40

Average = 30

minimize avg  
turnaround time!

# FIFO / FCFS



# FIFO / FCFS

**FIFO:** First In, First Out

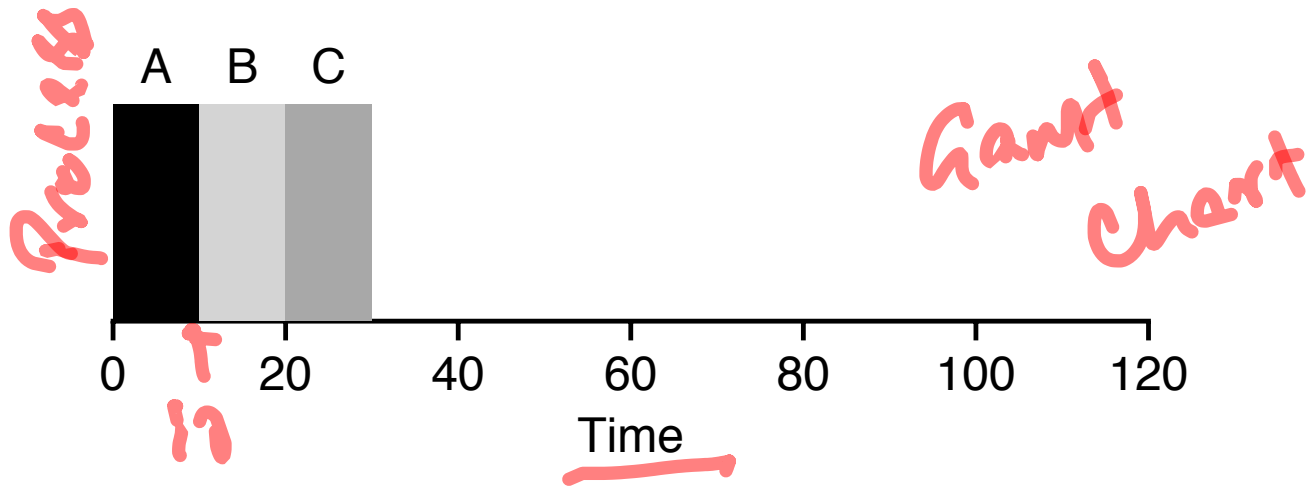
**FCFS:** First Come, First Served

Job	Arrival(s)	run time (s)
A	~0	10
B	~0	10
C	~0	10

# FIFO / FCFS

Job	Arrival(s)	run time (s)
A	~0	10
B	~0	10
C	~0	10

Average  
Turnaround  
Time ?



10, 20, 30  
= 20



# ASSUMPTIONS

- ~~1. Each job runs for the same amount of time~~
- 2. All jobs arrive at the same time
- 3. All jobs only use the CPU (no I/O)
- 4. Run-time of each job is known

# 2-MINUTE QUIZ

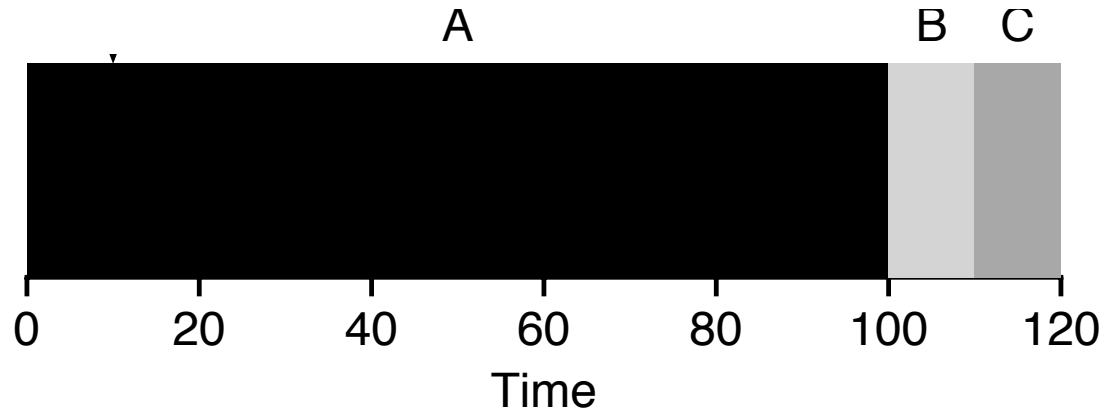
How will FIFO perform without this assumption ?

What scenarios can lead to bad performance?

Avg Turn around  
time

# BIG FIRST JOB

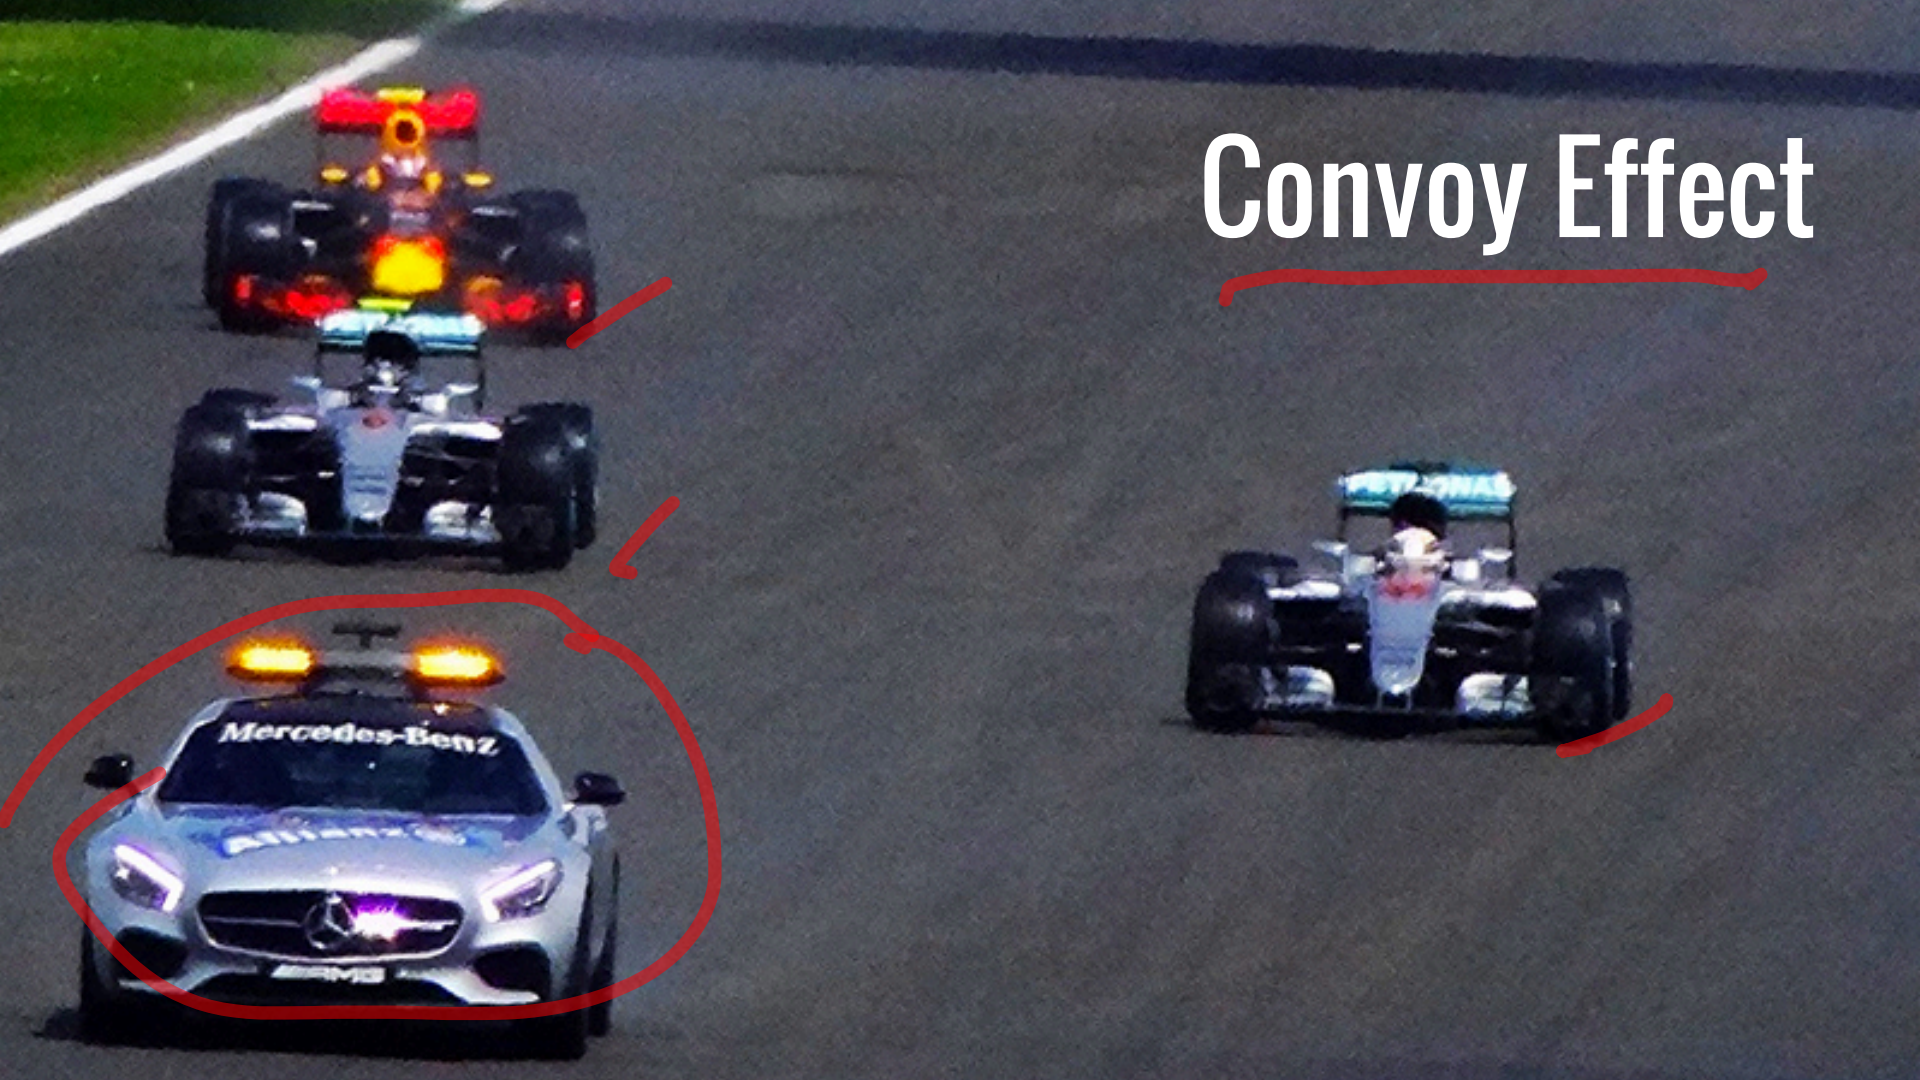
Job	Arrival(s)	run time (s)
A	~0	100
B	~0	10
C	~0	10



Average  
Turnaround  
Time

$$(100 + 110 + 120) / 3 = 110s$$

# Convoy Effect



# CHALLENGE

Turnaround time suffers when short jobs must wait for long jobs

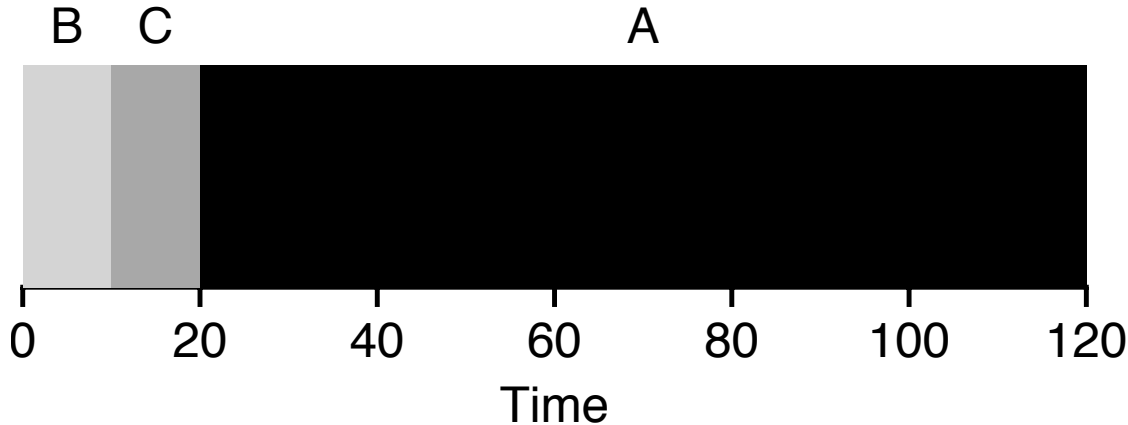
New scheduler:

SJF (Shortest Job First)

Choose job with smallest run\_time!

# SHORTEST JOB FIRST (SJF)

Job	Arrival(s)	run time (s)
A	~0	100
B	~0	10
C	~0	10



Average  
Turnaround  
Time

$$(10 + 20 + 120) / 3 = 50s!$$

FIFO: 110s ?!

# ASSUMPTIONS

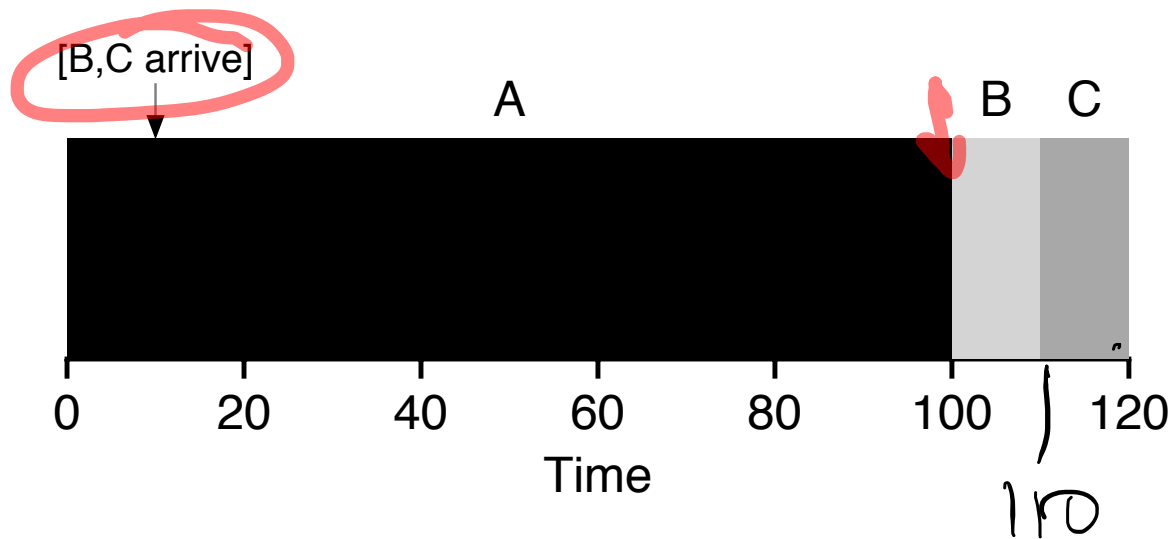
- ~~1. Each job runs for the same amount of time~~
- ~~2. All jobs arrive at the same time~~
- 3. All jobs only use the CPU (no I/O)
- 4. Run-time of each job is known

Job	Arrival(s)	run time (s)
A	~0	100
B	10	10
C	10	10

Average Turnaround Time with SJF?



Job	Arrival(s)	run time (s)
A	~0	100
B	10	10
C	10	10



Average  
Turnaround  
Time ?

$$(100 + 110 + 110) / 3$$

$$= 103.33$$

$$100 + 110 + 110$$

# PREEMPTIVE SCHEDULING

Prev schedulers:

FIFO and SJF are non-preemptive

Only schedule new job when previous job voluntarily relinquishes CPU

New scheduler:

Preemptive: Schedule different job by taking CPU away from running job

STCF (Shortest Time-to-Completion First)

Always run job that will complete the quickest

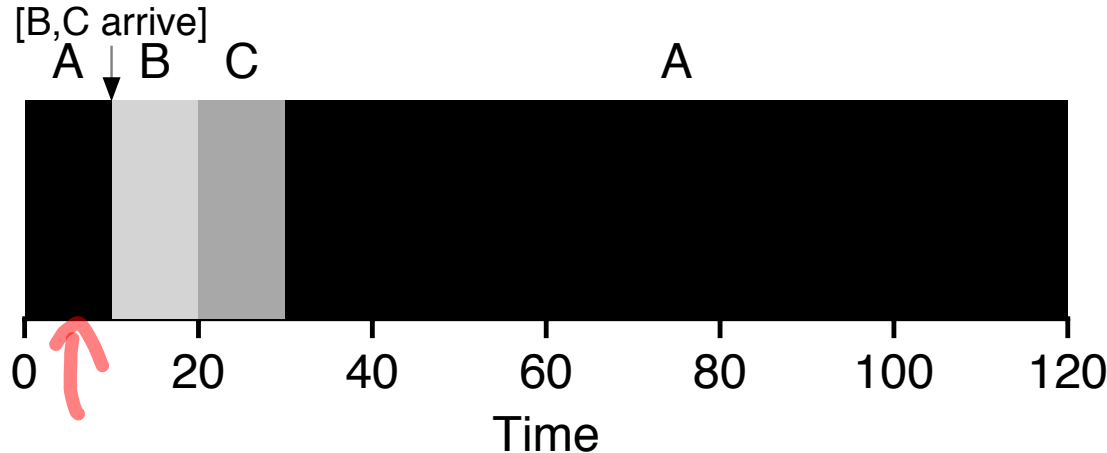
# PREMPTIVE SCTF

Job	Arrival(s)	run time (s)
A	~0	100
B	10	10
C	10	10

Handwritten calculations for Average Turnaround Time:

$$\frac{90 + 10 + 10}{3} = 36.67$$

Average  
Turnaround  
Time



$$(10 + 20 + 120) / 3 = 50s$$

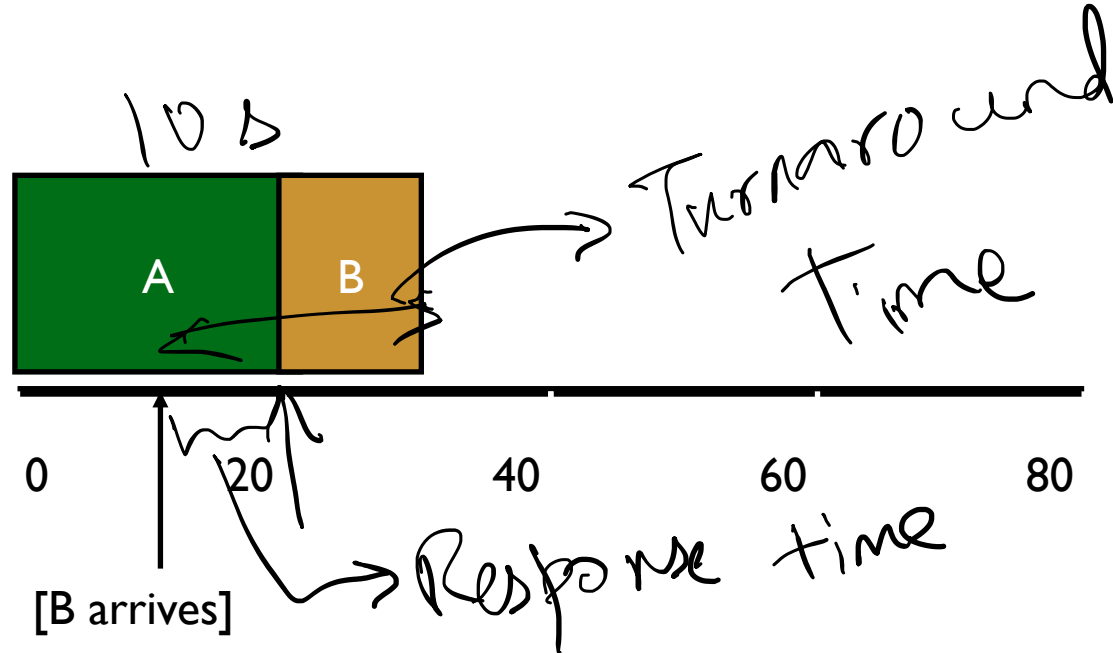
Handwritten red arrow pointing to 0 on the time axis and the number 100 below it.

# METRIC 2: RESPONSE TIME

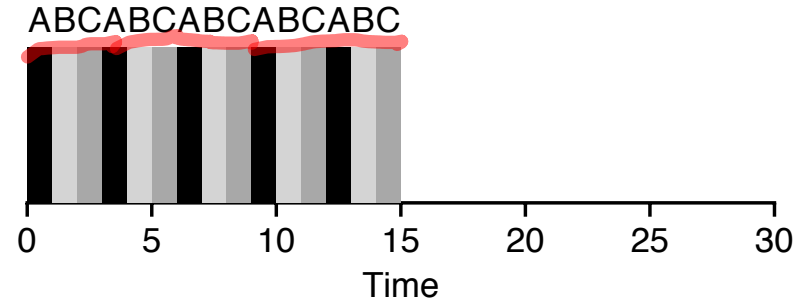
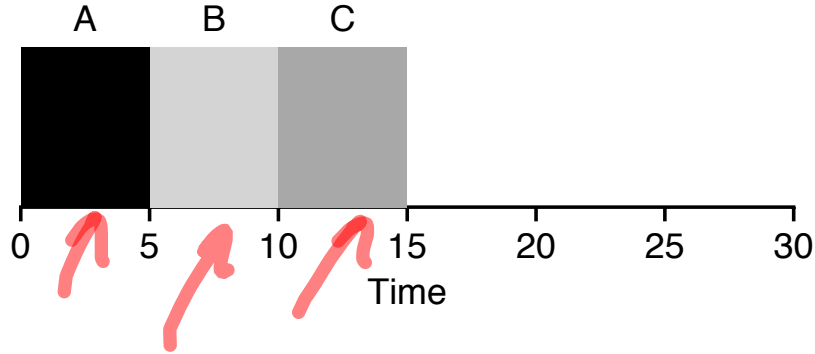
Response time = first\_run\_time - arrival\_time

B's turnaround: 20s

B's response: 10s



# ROUND ROBIN SCHEDULER

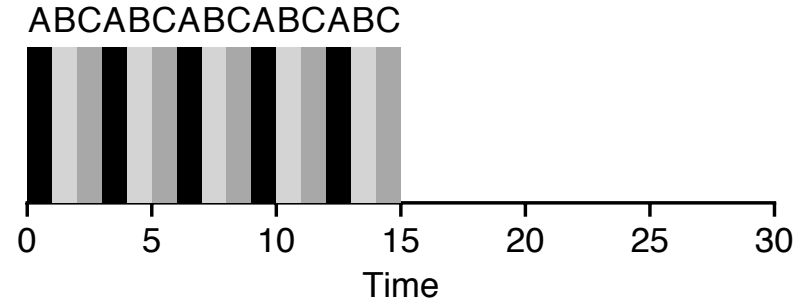
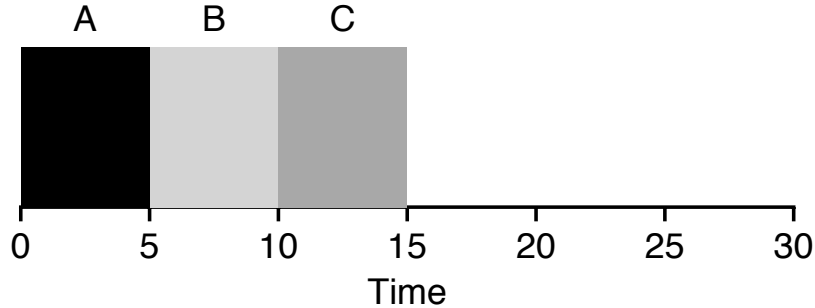


Average Response Time

$$(0 + 5 + 10)/3 = 5s$$

$$(0 + 1 + 2)/3 = 1s$$

# 2-MINUTE QUIZ



Av. 10s ?

14s

Turn  
around  
time

What is the turnaround time for two cases ?  
Is round robin better or worse?

# TRADE-OFFS

Round robin increases turnaround time decreases response time

Tuning challenges:

What is a good time slice for round robin?

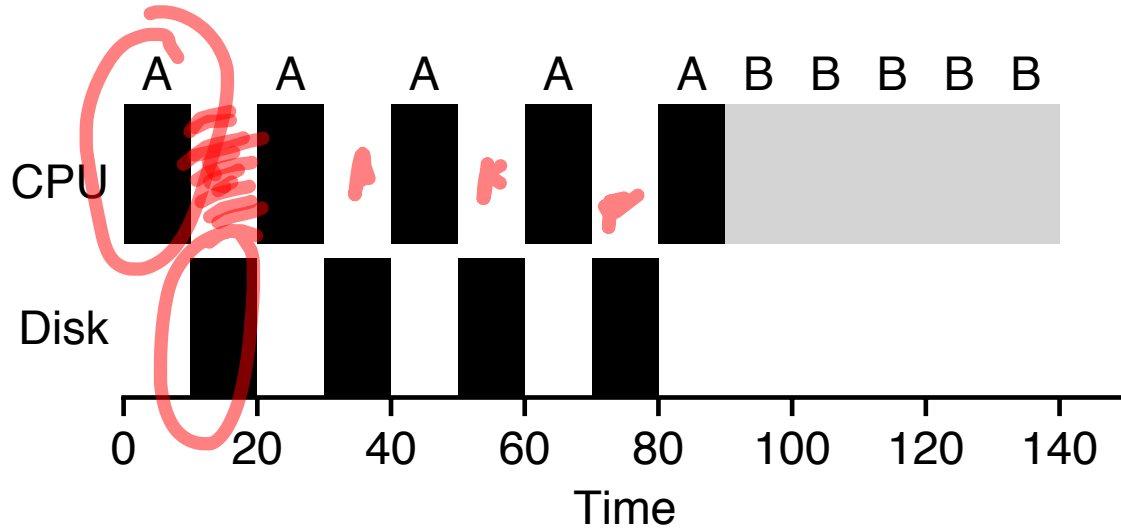
What is the overhead of context switching?

# ASSUMPTIONS

- ~~1. Each job runs for the same amount of time~~
- ~~2. All jobs arrive at the same time~~
- ~~3. All jobs only use the CPU (no I/O)~~
4. Run-time of each job is known

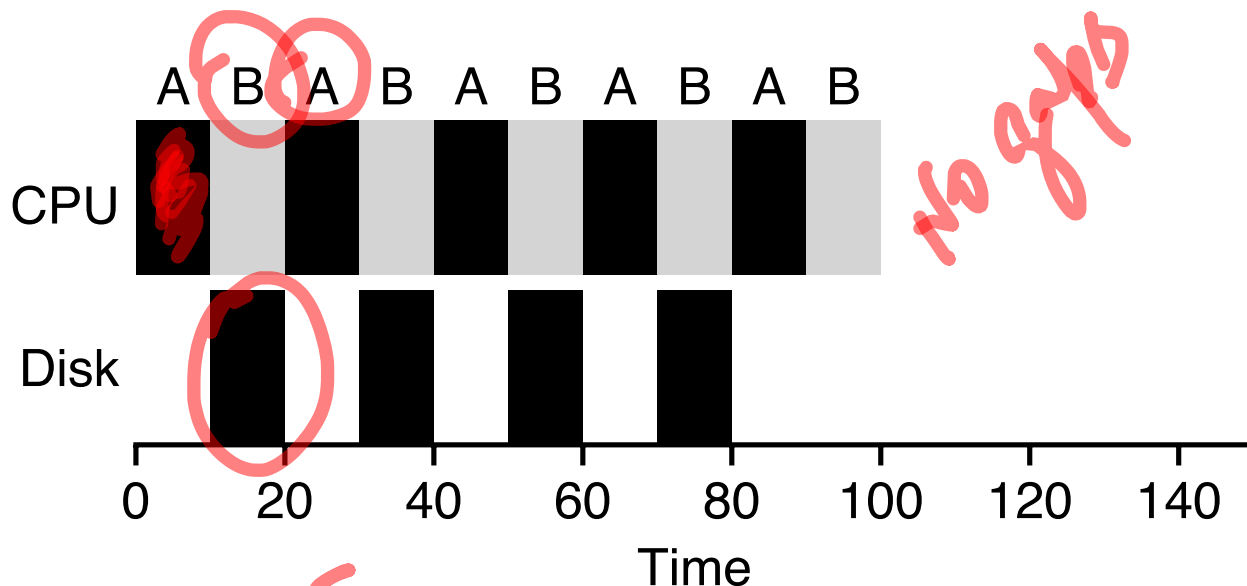


# NOT IO AWARE



Job holds on to CPU while blocked on disk!

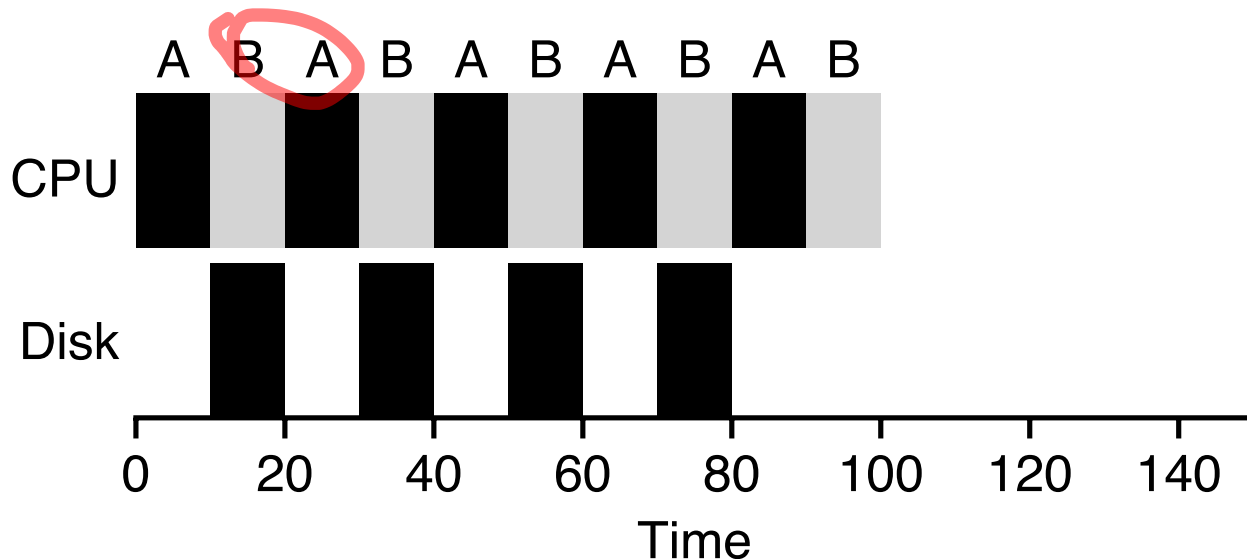
# I/O AWARE SCHEDULING



Treat Job A as 3 separate CPU bursts.

When Job A completes I/O, another Job A is ready

# I/O AWARE SCHEDULING



Each CPU burst is shorter than Job B

With SCTF, Job A preempts Job B

Treat Job A as 3 separate CPU bursts.

When Job A completes I/O, another Job A is ready

# ASSUMPTIONS

- ~~1. Each job runs for the same amount of time~~
- ~~2. All jobs arrive at the same time~~
- ~~3. All jobs only use the CPU (no I/O)~~
4. Run time of each job is known

# MULTI-LEVEL FEEDBACK QUEUE

---

MLFQ

# MLFQ: GENERAL PURPOSE SCHEDULER

Must support two job types with distinct goals

- “interactive” programs care about response time
- “batch” programs care about turnaround time

Key-1

Zip

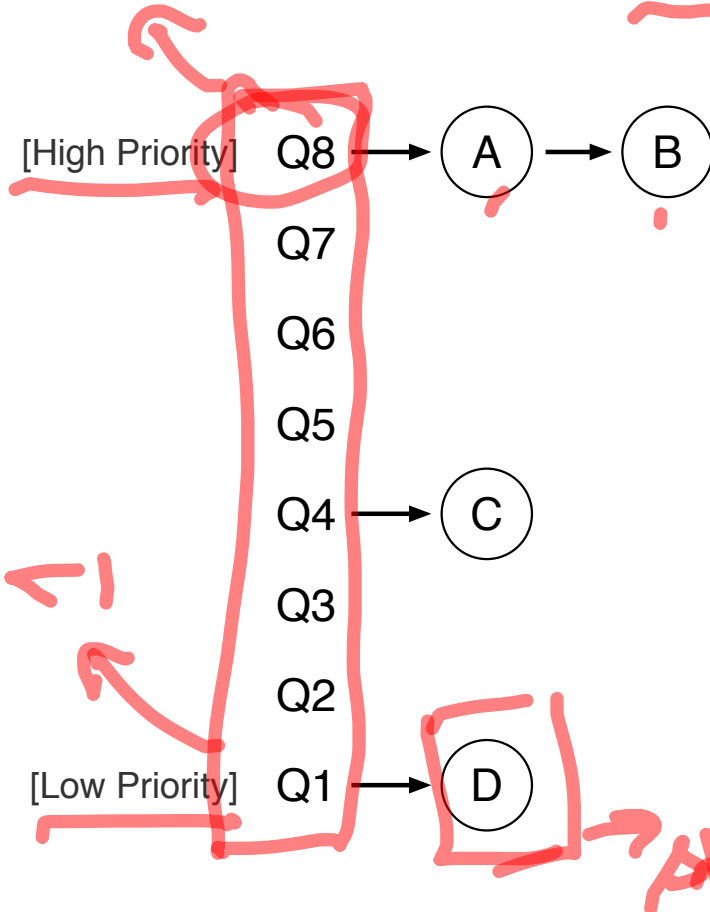
Approach:

Multiple levels of round-robin

Each level has higher priority than lower level

Can preempt them

# MLFQ EXAMPLE



“Multi-level” – Each level is a queue!

Rules for MLFQ

Rule 1: If  $\text{priority}(A) > \text{Priority}(B)$   
A runs

Rule 2: If  $\text{priority}(A) == \text{Priority}(B)$ ,  
A & B run in RR

# CHALLENGES

How to set priority?

What do we do when a new process arrives?

Does a process stay in one queue or move between queues?

Approach: Use past behavior of process to predict future!

Guess how CPU burst (job) will behave based on past CPU bursts



# MORE MLFQ RULES

Rule 1: If  $\text{priority}(A) > \text{Priority}(B)$ , A runs

Rule 2: If  $\text{priority}(A) == \text{Priority}(B)$ , A & B run in RR

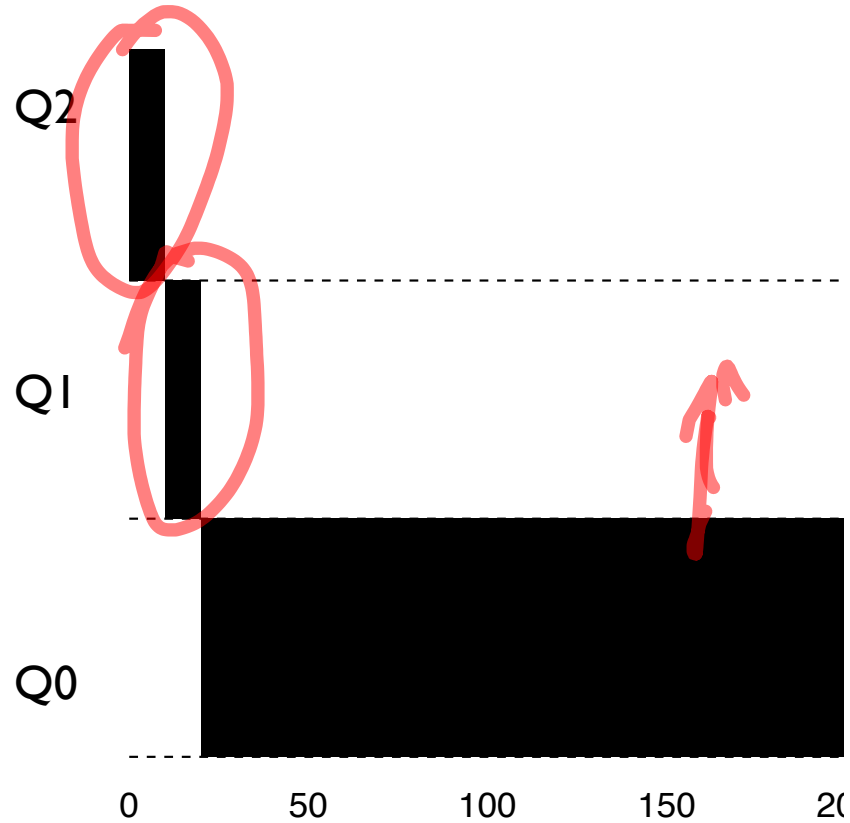
Rule 3: Processes start at top priority

Rule 4: If job uses whole slice, demote process

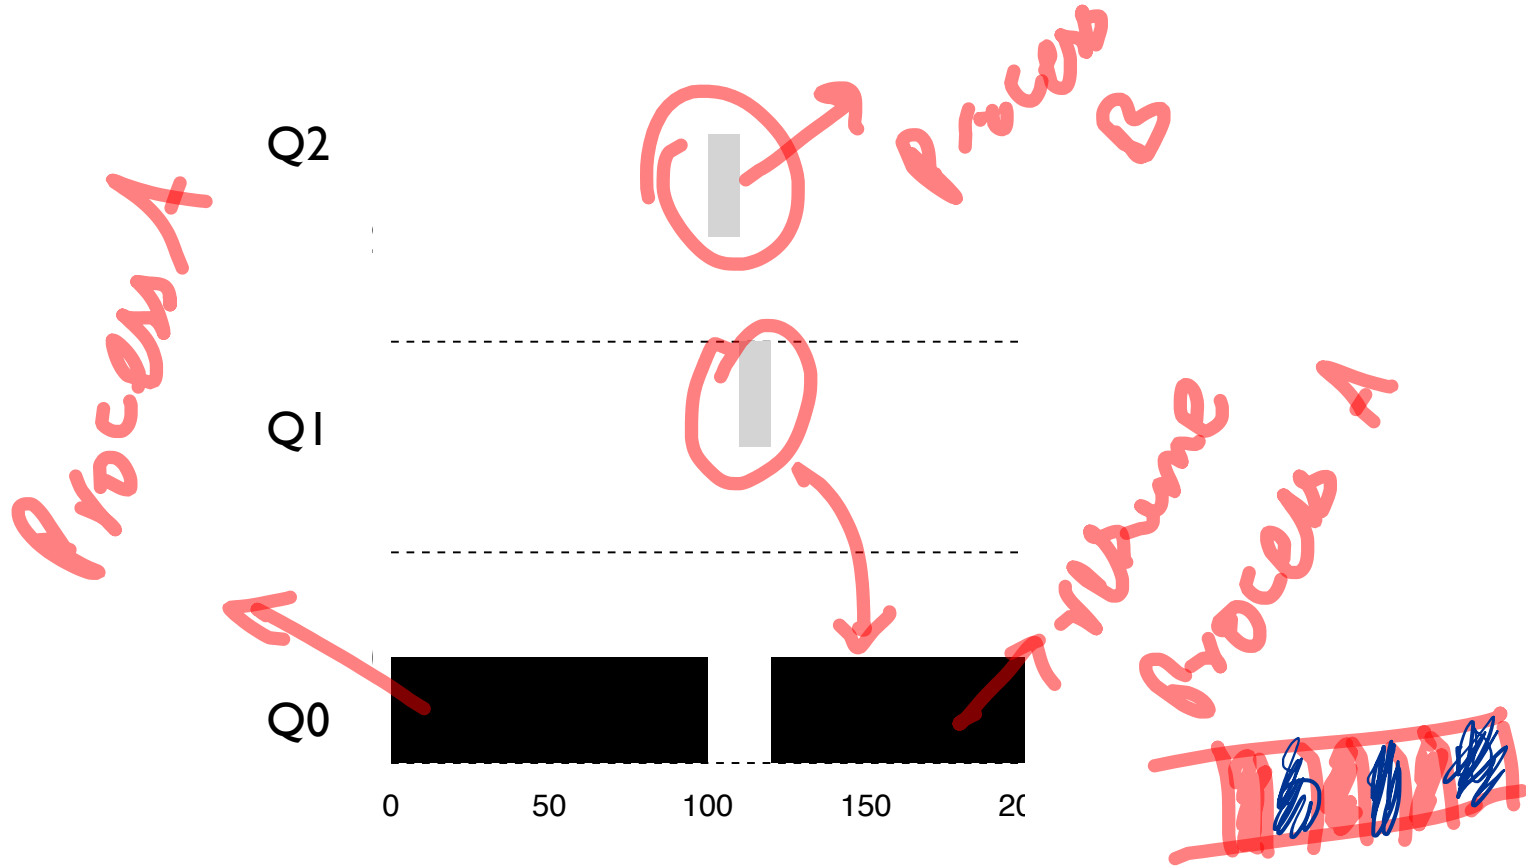
(longer time slices at lower priorities)

//

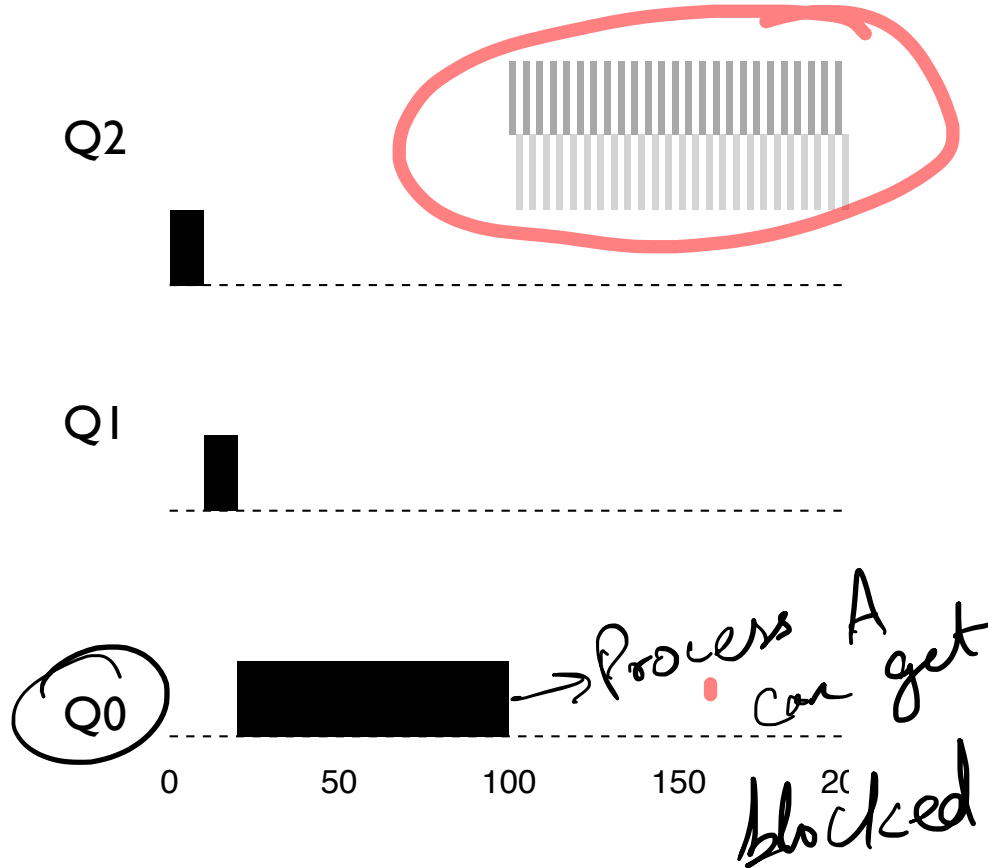
# ONE LONG JOB



# INTERACTIVE PROCESS JOINS



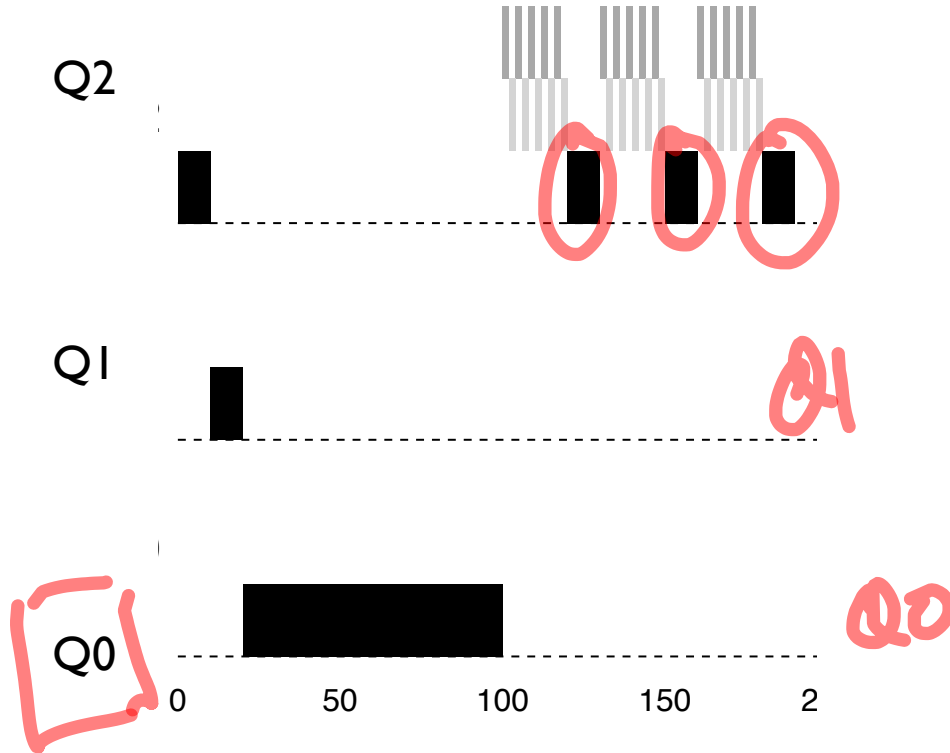
# MLFQ PROBLEMS?



What is the problem with this schedule ?

starvation

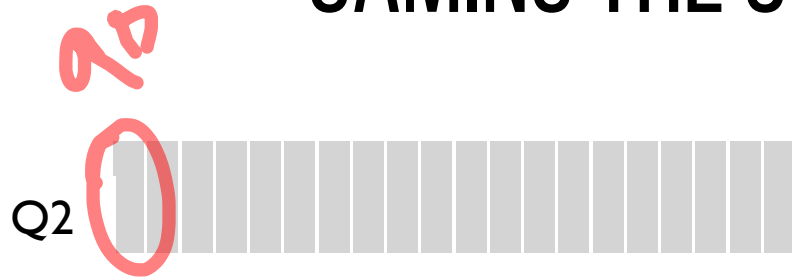
# AVOIDING STARVATION



Problem: Low priority job may never get scheduled

Periodically **boost** priority of all jobs (or all jobs that haven't been scheduled)

# GAMING THE SCHEDULER ?



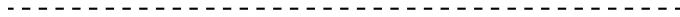
9s 1s



timer interrupt  
time slice 10s

Job could trick scheduler by doing I/O just before time-slice end

Q1



Q0



0 50 100 150 200

Account for **total run time** at priority  
Downgrade when exceed threshold

# SUMMARY

## Scheduling Policies

Understand workload characteristics like arrival, CPU, I/O

Scope out goals, metrics (turnaround time, response time)

## Approach

Trade-offs based on goals, metrics (RR vs. SCTF)

Past behavior is good predictor of future behavior?

MLFQ

# NEXT STEPS

Project 1a: Due Jan 31 (Thursday) at 11.59pm

Project 1b: Out on Jan 30<sup>th</sup>

Thursday class, discussion

- More advanced scheduling policies

- Summary / review of process, CPU scheduling

- xv6 introduction, walk through

- Go through xv6 context switch / syscall?